

**Report Title:** End-to-End Deep Learning Architecture for Multivariate Energy  
Forecasting

**Author:** Yahal Sineth Yahampath

**Date:** 11<sup>th</sup> June 2026

## Contents

|                                                      |    |
|------------------------------------------------------|----|
| 1. Introduction .....                                | 3  |
| 2. Data Insights .....                               | 3  |
| 2.1 Cyclicity and Diurnal Patterns .....             | 3  |
| 2.2 Volatility and Correlation .....                 | 4  |
| 2.3 Consumption Distribution and Anomalies .....     | 5  |
| 3. Preprocessing .....                               | 6  |
| 4. Feature Engineering.....                          | 6  |
| 5. Model Design .....                                | 8  |
| 5.1 Deep Learning Architecture (CNN-LSTM).....       | 8  |
| 5.2 Network Configuration and Hyperparameters.....   | 8  |
| 6. Results.....                                      | 9  |
| 6.1 Evaluation Metrics Showdown .....                | 9  |
| 6.2 Architectural Insights & Algorithm Failure ..... | 10 |
| 7. Model Optimization .....                          | 13 |
| 7.1 Hyperparameter Search Space .....                | 13 |
| 7.2 Optimization Outcomes .....                      | 13 |
| 8. Challenges and Solutions .....                    | 14 |
| 9. Conclusion & Future Work .....                    | 14 |
| 10. References .....                                 | 16 |

# 1. Introduction

The objective of this project was to engineer a robust, production-grade Machine Learning pipeline capable of predicting multivariate time-series energy consumption. Accurately forecasting energy demand is critical for operational load balancing and grid stability. This pipeline establishes benchmark performance using traditional machine learning algorithms (Linear Regression, Random Forest, XGBoost) and contrasts them against a dynamically optimized Deep Learning architecture (CNN-LSTM) to determine the most effective forecasting strategy.

## 2. Data Insights

The foundational phase of the pipeline involved rigorous exploratory data analysis (EDA) of the raw temporal data to identify underlying consumption patterns.

### 2.1 Cyclicity and Diurnal Patterns

Visual analysis of the time-series data revealed distinct daily and seasonal energy consumption cycles. Energy spikes are highly correlated with specific hours of the day, indicating peak operational hours or human behavioral patterns.





Figure 2.1: Time-series visualization illustrating diurnal energy consumption cycles.

## 2.2 Volatility and Correlation

The data exhibited sudden, massive consumption spikes. While traditional linear observation suggested these spikes were chaotic, correlation mapping indicated they were tied to specific meteorological conditions and temporal features.




**Figure 2.2:** Correlation heatmap identifying the relationships between raw weather variables, time, and energy consumption.

**Key Observation:** The exploratory analysis demonstrated that predictive models cannot effectively interpret raw, linear time. Because energy consumption tightly correlates with specific times of the day, representing time as a continuous mathematical cycle became essential for accurate forecasting (detailed in Section 4).

## 2.3 Consumption Distribution and Anomalies

Analysis of the data distribution highlighted severe positive skewness due to massive, infrequent energy spikes.



**Figure 2.3:** Distribution of energy consumption highlighting extreme load anomalies.

**Key Observation:** The presence of these extreme values mathematically justified the implementation of robust feature scaling (Section 3) and dictated that Root Mean Squared Error (RMSE) be heavily weighted during model evaluation, as it penalizes large predictive misses during peak loads.

### 3. Preprocessing

To ensure the data was scientifically valid before entering the feature engineering phase, strict preprocessing protocols were executed:

- **Cleaning & Missing Values:** The dataset was scanned for null values and temporal gaps. Forward-fill imputation was prioritized for minor gaps to maintain the chronological integrity of the time series without injecting synthetic averages.
- **Scaling:** Because neural networks are highly sensitive to unscaled variance (e.g., comparing a temperature of 85°F to an energy consumption metric of 10,000 watts), the continuous features were normalized using standard scaling techniques. This ensured that no single variable dominated the loss gradient during gradient descent.
- **Temporal Data Isolation:** To prevent data leakage a fatal flaw in time-series forecasting standard randomized cross-validation was abandoned. Instead, a strict sequential 70/20/10 split (Train/Validation/Test) was enforced. The Test set was strictly isolated and only accessed during the final evaluation phase.

### 4. Feature Engineering

To translate human-readable observations into machine-readable continuous variables, the pipeline dynamically generated physics- and cycle-based features.

- **Cyclical Time Encoding:** hour\_sin, hour\_cos, month\_sin, and month\_cos features were created. This trigonometric transformation allows the models to mathematically understand the continuous loop of time (i.e., understanding that 23:00 and 00:00 are chronologically adjacent, not distant extremes).
- **Domain Metrics:** Physics-based features, specifically temperature differentials (Delta\_T), were engineered to directly capture the environmental drivers of HVAC energy consumption.
- **Algorithmic Feature Selection:** To prevent the "curse of dimensionality," a diagnostic Random Forest Regressor was utilized to evaluate the predictive power of every feature.




Figure 4: Algorithmic Feature Importance highlighting the mathematical weighting of engineered variables.

Features demonstrating minimal predictive variance based on Gini impurity reduction metrics were systematically dropped, ensuring the final models learned purely from high-signal data.

## 5. Model Design

The forecasting architecture was designed as a comparative showdown between traditional Machine Learning baselines (Linear Regression, Random Forest, XGBoost) and a custom Deep Learning network.

### 5.1 Deep Learning Architecture (CNN-LSTM)

For the advanced neural network approach, a hybrid CNN-LSTM (Convolutional Neural Network - Long Short-Term Memory) architecture was engineered.

- **CNN Integration (Conv1D):** Standard LSTMs can struggle with sudden, high-frequency local noise. A 1D Convolutional layer was implemented at the input to act as a feature extractor, smoothing out micro-volatility and passing clean, local temporal patterns to the recurrent layers.
- **LSTM Integration:** The LSTM layer was utilized to capture long-term sequential dependencies and maintain an internal memory state of the diurnal and seasonal cycles over time.

### 5.2 Network Configuration and Hyperparameters

- **Activation Functions:** The ReLU (Rectified Linear Unit) activation function was utilized across dense and convolutional layers to prevent the vanishing gradient problem and accelerate convergence.
- **Optimizer:** The Adam optimizer was selected for its adaptive learning rate capabilities, adjusting the gradient descent step size dynamically based on the geometry of the loss surface.
- **Loss Function:** Mean Squared Error (MSE) was set as the primary loss function during training to heavily penalize the model for missing massive energy spikes (as opposed to Mean Absolute Error, which treats all errors linearly).




## 6. Results

The isolated Test set was evaluated against all serialized models. The performance metrics completely disrupted initial assumptions regarding algorithm superiority.

### 6.1 Evaluation Metrics Showdown

| Model             | Mean Absolute Error | Root Mean Squared Error |
|-------------------|---------------------|-------------------------|
| Linear Regression | 0.7759              | 1.3039                  |
| Tuned CNN-LSTM    | 1.0251              | 1.8939                  |
| Random Forest     | 1.1835              | 1.8225                  |
| XGBoost           | 0.5674              | 1.2617                  |

Table 1: Metrics comparison between the Models



**Figure 6.1:** Metric comparison highlighting the superior predictive accuracy of the XGBoost algorithm.

## 6.2 Architectural Insights & Algorithm Failure

The metrics reveal a critical lesson in tabular time-series forecasting: Deep Learning is not a silver bullet.

While the CNN-LSTM was engineered with a 24-step sequence window to capture long-term temporal dependencies, it was fundamentally outperformed by the gradient-boosted XGBoost model. Neural networks typically require massive, high-dimensional datasets to reach their full potential. In contrast, XGBoost armed with rigorous hyperparameter tuning (optimized learning rates and tree depth) and the implementation of RobustScaler to handle extreme outliers, was highly efficient at mapping the non-linear cyclical engineered features to energy consumption.

Interestingly, the traditional Linear Regression model secured second place. This is a direct testament to the rigorous mathematical feature engineering executed in Phase 1. Because the cyclical nature of time was perfectly encoded into sine and cosine waves, the Linear Regression model easily mapped those waves to energy consumption without requiring complex, non-linear architecture.







Figure 6.2: Time-series prediction tracking. Note how XGBoost aggressively and successfully maps the extreme energy spikes, while the CNN-LSTM suffers from mean-collapse, predicting a flat baseline.



**Figure 6.3:** Actual vs. Predicted 45-degree scatter reference. The horizontal clustering of the CNN-LSTM predictions reveals a systematic failure to predict high-consumption events.



**Figure 6.4:** Residual scatter plot. The perfect diagonal trajectory indicates that as actual energy consumption scales up, the model's predictive error scales at a 1:1 ratio, proving it is entirely blind to peak anomalies.



**Figure 6.5:** Error distribution. The heavy right-skew visually quantifies the systematic underprediction of extreme load events.

## 7. Model Optimization

Rather than manually guessing network parameters, the pipeline utilized KerasTuner to algorithmically hunt for the optimal Deep Learning architecture.

### 7.1 Hyperparameter Search Space

The tuner dynamically adjusted the number of Conv1D filters, LSTM units, learning rates, and Dropout layers over multiple trial epochs, validating against the 20% validation set to prevent data leakage.

### 7.2 Optimization Outcomes

The most significant finding from the tuning phase was diagnosing the CNN-LSTM's structural failure, known as Mean Collapse. The heavy-tailed distribution of the energy spikes heavily penalized the network during gradient descent. To minimize overall MSE, the KerasTuner aggressively optimized for safety, selecting a severe Dropout rate (0.4).

This over-regularization caused the neural network to completely ignore peak anomalies and continuously predict the median energy consumption. Conversely, the 2D tree-based models (XGBoost) successfully utilized manually engineered lag features (`Appliances_lag_10m`) to perform highly accurate autoregression. This proves that for this specific dataset, explicit feature engineering combined with gradient boosting is vastly superior to implicit sequence learning via LSTMs.

## 8. Challenges and Solutions

Developing a production-grade time-series forecasting pipeline presents unique mathematical and architectural hurdles. The following critical challenges were identified and systematically resolved:

- **Challenge 1: Temporal Data Leakage**
  - *Issue:* Standard Machine Learning cross-validation randomly shuffles data. If future energy data is accidentally shuffled into the training set, the model "cheats" by seeing the future, resulting in artificially high accuracy during training but catastrophic failure in production.
  - *Solution:* Standard cross-validation was entirely abandoned. The `build_dataset.py` orchestrator was programmed to enforce a strict, chronological 70/20/10 sequence, guaranteeing the Test set remained strictly isolated as a true representation of the unseen future.
- **Challenge 2: Extreme Load Volatility**
  - *Issue:* The energy grid experiences sudden, massive consumption spikes that can destabilize model weights during gradient descent.
  - *Solution:* A dual approach was utilized. First, a 1D Convolutional layer was integrated into the Deep Learning architecture to mathematically smooth local temporal noise. Second, Root Mean Squared Error (RMSE) was prioritized over standard MAE to heavily penalize the model for missing large consumption anomalies.
- **Challenge 3: Scaling for Heavy-Tailed Anomalies & Baseline Parity**
  - *Issue:* Initial evaluations utilized standard `MinMaxScaler` algorithms which collapsed under the weight of massive, unpredictable energy spikes. Furthermore, untuned tree-based baselines artificially skewed the comparative analysis.
  - *Solution:* Deployed `RobustScaler` leveraging Interquartile Ranges to safely process extreme load anomalies without clipping data. Subsequently implemented `RandomizedSearchCV` for hyperparameter tuning across XGBoost and Random Forest. This parity protocol resulted in XGBoost reducing MAE to an industry-leading 0.5674, proving the superiority of tuned gradient boosting for this specific tabular architecture.

## 9. Conclusion & Future Work

This project successfully engineered a fully automated, end-to-end pipeline for multivariate energy forecasting. By strictly enforcing the Single Responsibility Principle, isolating temporal data, and implementing robust scaling, the architecture is highly reproducible and production-ready.

The comparative analysis yielded a critical engineering insight: complex architectures do not automatically equate to superior performance. The champion model was XGBoost, which leveraged the engineered cyclical features to achieve an MAE of 0.5674, significantly outperforming the deep learning CNN-LSTM network. This outcome underscores the principle that for tabular, highly-structured temporal data, well-tuned gradient-boosted trees remain an industry gold standard.

### Areas for Future Work:

- **Real-Time Integration:** Expanding the data ingestion module to pull live meteorological data via APIs to adjust forecasts in real-time.
- **Cloud Deployment & Orchestration:** Containerizing the training and inference modules using Docker and deploying the pipeline as scalable microservices (potentially utilizing Azure infrastructure) to enable continuous automated retraining.
- **Extended Deep Learning Context:** While the CNN-LSTM underperformed the tree models on this specific dataset volume, expanding the training data and investigating the efficacy of temporal attention mechanisms (e.g., TimeSformer) could unlock superior performance over longer temporal horizons.

## 10. References

1. **Chollet, F. et al. (2015).** *Keras: The Python Deep Learning library*. Keras.io.
2. **Pedregosa, F. et al. (2011).** *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12, 2825-2830.
3. **Chen, T., & Guestrin, C. (2016).** *XGBoost: A Scalable Tree Boosting System*. Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining.
4. **O'Malley, T. et al. (2019).** *KerasTuner: An easy-to-use, scalable hyperparameter optimization framework*. Keras.io.
5. **McKinney, W. (2010).** *Data Structures for Statistical Computing in Python*. Proceedings of the 9th Python in Science Conference.